

final_project_knitted

August 6, 2022

```
[ ]: %reset -f

[ ]: # Import stuff
# !pip install pandas numpy seaborn sklearn nltk missingno tensorflow dask
    ↪hvplot
# !pip install xgboost --upgrade
import os
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
import plotly.express as px
from plotly.subplots import make_subplots
import plotly.graph_objects as go
import missingno as msno
from sklearn import model_selection, preprocessing, ensemble, neighbors,
    ↪linear_model, svm, metrics
from nltk.util import ngrams
from imblearn import over_sampling, under_sampling, pipeline
from PIL import Image
from sklearn.experimental import enable_iterative_imputer
from sklearn.impute import IterativeImputer, SimpleImputer
from sklearn import metrics
from xgboost import XGBClassifier, XGBRegressor # 1.6.1
from sklearn.pipeline import Pipeline
import itertools
from sklearn.utils import class_weight
import plotly.io as pio
pio.renderers.default = "notebook+pdf"
```

First we'll define some useful functions for later:

```
[ ]: # Define useful functions

"""
Returns dictionary of textual column values and their frequency in dataframe
"""
def get_column_values(df, exceptions=[]):
```

```

n = len(df)
text_columns = df.select_dtypes(exclude=[np.number])
column_values = {}
for column in text_columns:
    # Make frequencies dictionary and calculated percent of missing values
    if column not in exceptions:
        column_values[column] = df[column].value_counts().to_frame().
↪reset_index()
        column_values[column].columns = ['value', 'count']
        column_values[column]['pc'] = column_values[column]['count'] / n * 100
↪100
    return column_values

"""
Cleans text values in dataframe for EDA purposes:
0. Replaces spaces with underscore
1. Removes non-ASCII characters
2. Transforms to lowercase
3. Fills null values with "NA"
4. Removes punctuation characters (except underscore and .)
"""
from string import punctuation
def clean_text_values(df):
    # punc = [p for p in punctuation if p not in ['_']]
    column_values = get_column_values(df)
    for column in column_values:
        df[column].str.encode('ascii', 'ignore').str.decode('ascii')
        df[column] = df[column].str.lower()
        df[column].fillna('NA', inplace=True)
        if column != 'category':
            for p in punctuation:
                df[column] = df[column].str.replace(p, '', regex=False)
            df[column] = df[column].str.replace(r'\s+', ' ', regex=True)
    return df

"""
Recieves df, categories and column name as arguments and returns a Histogram
↪object of that column divided by the different categories, without
zeros and values above the 99th percentile to clear extreme values which mess
↪up the histogram (calculated after removing zeros).
"""
def create_column_histogram_by_category(df, categories, column):
    all_data = []
    num_99_percentile = 0
    for category in categories:

```

```

        data_no_zeros = df[(df['category'] == category) & (df[column] > 0)]
        data = data_no_zeros[(data_no_zeros[column] <= data_no_zeros[column].
↪quantile(0.99))]
        num_99_percentile = len(data_no_zeros) - len(data)
        all_data.append(go.Histogram(
            x=data[column],
            name=category,
            opacity=0.7
            # ,xbins={
            #     'start': 0,
            #     'end': data.max(),
            #     'size': data.max() / 500
            # }
        ))
    )
    num_zeros = len(df[df[column] == 0])
    layout = go.Layout(
        barmode='overlay',
        title='{ } histograms divided by class | No. of zeros: { } | No.
↪above 99th percentile: { }\n| Percent "bad" rows: { }%'.format(
            column.capitalize().replace('_', ' '), num_zeros,
↪num_99_percentile, round((num_zeros + num_99_percentile) / len(df) * 100, 3))
    )
    fig = go.Figure(data=all_data, layout=layout)
    return fig

def add_ngrams_columns(df, column, n, return_column=False):
    new_ngrams_column = '{ }_{ }gram'.format(column, n)
    df[new_ngrams_column] = df[df[column].notna()][column].apply(lambda row:
↪list(ngrams(row.split(' '), n)))
    if return_column:
        return df, new_ngrams_column
    return df

def get_ngrams(df, column, n, remove_column=False):
    df, new_ngrams_column = add_ngrams_columns(df, column, n,
↪return_column=True)
    categories = list(df['category'].unique())
    n_grams = {cat: { } for cat in categories}
    for category in categories:
        cat_df = df[df['category'] == category]
        for ngram_list in cat_df[new_ngrams_column].
↪notna()[new_ngrams_column]:
            if type(ngram_list) == float:

```

```

        print(ngram_list)
    for n_gram in ngram_list:
        if n_gram in n_grams[category].keys():
            n_grams[category][n_gram] += 1
        else:
            n_grams[category][n_gram] = 1
    if remove_column:
        df = df.drop([new_ngrams_column], axis=1)
    return df, n_grams

def get_ngrams_top_k(df, column, n, k=1, remove_column=False):
    df, column_ngrams = get_ngrams(df, column, n, remove_column)
    categories = list(df['category'].unique())
    top_20_column_ngrams_by_category = {}
    for category in categories:
        top_column_ngrams = sorted(column_ngrams[category],
    ↪key=column_ngrams[category].get, reverse=True)[:k]
        top_20_column_ngrams = [{" ".join(key), column_ngrams[category][key]}]
    ↪for key in column_ngrams[category].keys() if key in top_column_ngrams
        top_20_column_ngrams = pd.DataFrame(top_20_column_ngrams,
    ↪columns=['ngram', 'count']).sort_values('count')
        top_20_column_ngrams_by_category[category] = top_20_column_ngrams
    return df, top_20_column_ngrams_by_category

def get_dims(file):
    img = Image.open(file)
    try:
        h, w, d = np.array(img).shape
    except:
        h, w = np.array(img).shape
    img.close()
    return h, w

```

Let us load the data:

```
[ ]: dataset = pd.read_csv('data/food_train.csv')
fact_nutrients = pd.read_csv('data/food_nutrients.csv')
dim_nutrients = pd.read_csv('data/nutrients.csv')
```

Then merge and clean it a bit:

```
[ ]: # Join nutrients tables and clean it
nutrients = pd.merge(fact_nutrients, dim_nutrients, how='left',
    ↪on='nutrient_id')
nutrients = clean_text_values(nutrients)
```

```

nutrients['name'] = nutrients['name'].str.replace(' ', '_')
nutrients['name'] = nutrients['name'] + '_' + nutrients['unit_name']
nutrients_column_values = get_column_values(nutrients)

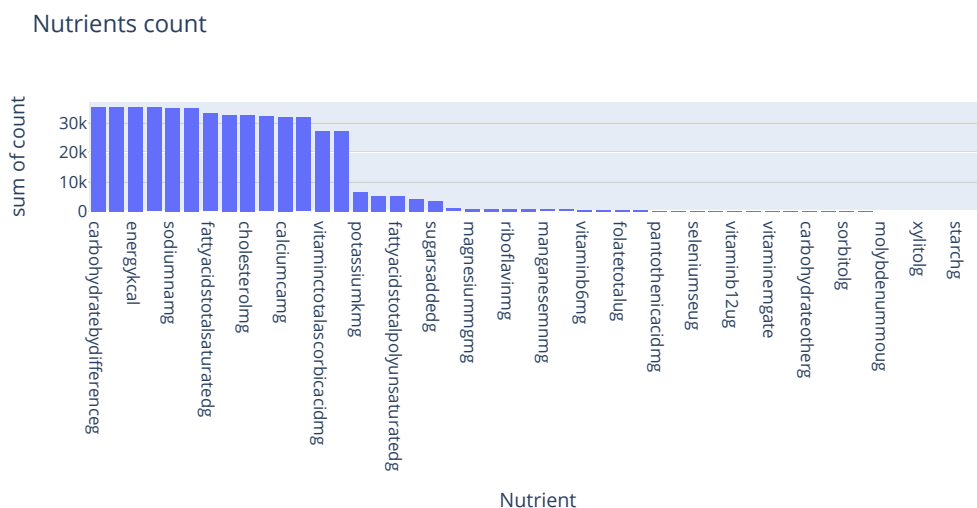
```

See how many records we have of each nutrient:

```

[ ]: # Draw histogram of nutrients using plotly
num_nutrients = nutrients.groupby('name')['idx'].count().reset_index()
num_nutrients.columns = ['nutrient', 'count']
num_nutrients['pc'] = num_nutrients['count'] / len(dataset) * 100
num_nutrients = num_nutrients[num_nutrients['count'] > 0]
num_nutrients = clean_text_values(num_nutrients)
px.histogram(
    num_nutrients.sort_values('count', ascending=False), x='nutrient',
    y='count', title='Nutrients count', width=800, height=400,
    labels={'nutrient': 'Nutrient'}
)

```



Now let's left join the nutrients to our snacks dataset:

```

[ ]: # Left join nutrients to dataset
joined_dataset = pd.merge(dataset, nutrients, how='left', on='idx').
    drop(['nutrient_id', 'unit_name'], axis=1)

```

Good ol' pivot table and join:

```

[ ]: # Pivot table - started by resetting index and rearranging columns but decided
    to join again to avoid pandas index trouble
pivoted_dataset = pd.pivot_table(

```

```

        joined_dataset,
        values='amount',
        index='idx',
        columns='name')

dataset_w_nutrients = pd.merge(dataset, pivoted_dataset, how='left', on='idx')

```

Change categories names and encode them for convenience:

```

[ ]: dataset_w_nutrients.loc[dataset_w_nutrients['category'] == 'cakes_cupcakes_snack_cakes', 'category'] = 'cakes'
      ↪ 'chips_pretzels_snacks', 'category'] = 'snacks'
dataset_w_nutrients.loc[dataset_w_nutrients['category'] == 'cookies_biscuits', 'category'] = 'cookies'
dataset_w_nutrients.loc[dataset_w_nutrients['category'] == 'popcorn_peanuts_seeds_related_snacks', 'category'] = 'seeds'
le = preprocessing.LabelEncoder()
le.fit(dataset_w_nutrients['category'])
dataset_w_nutrients['category_enc'] = le.transform(dataset_w_nutrients['category'])
le.classes_

```

```

[ ]: array(['cakes', 'candy', 'chocolate', 'cookies', 'seeds', 'snacks'],
          dtype=object)

```

Now we can split our data for initial analysis!

```

[ ]: # Split to perform EDA only on train data so conclusions won't later leak to test set
      ↪ test set
X = dataset_w_nutrients.drop(['category'], axis=1)
y = dataset_w_nutrients['category']
X_train, X_test, y_train, y_test = model_selection.train_test_split(X, y,
      ↪ test_size=0.2, random_state=1337, stratify=y)
eda_df = pd.concat([X_train, y_train], axis=1)

```

Clean text values:

```

[ ]: eda_df = clean_text_values(eda_df)

```

Hajime!


First of all, let's check for nulls:

```

[ ]: nulls = eda_df.isnull().sum().to_frame().reset_index()
      nulls.columns = ['column', 'nulls']
      nulls = nulls.sort_values('nulls')

```

```
px.bar(nulls, x='nulls', y='column', orientation='h', title='Nulls count',
      width=700, height=350)
#nulls.head(2)
```



```
[ ]: # Code to remove high nulls nutrients - would maybe use later

# eda_n = len(eda_df)
# nutrients_to_drop = []
# for nutrient in nutrients_column_values['name']['value']:
#     num_zeros = (eda_df[nutrient].isna()).astype(int).sum()
#     pc_zeros = num_zeros / eda_n * 100
#     if pc_zeros >= 80:
#         nutrients_to_drop.append(nutrient)
# eda_df = eda_df.drop(nutrients_to_drop, axis=1)
```


```
[ ]: eda_df, top20_ingredients_words_by_category = get_ngrams_top_k(eda_df,
      'ingredients', 1, 20)
```

Let's have a look at some of the common words in the ingredients column:


```
[ ]: # Get an idea of distinct words in 'ingredients' column
eda_df, top20_ingredients_words_by_category = get_ngrams_top_k(eda_df,
      'ingredients', 1, 20)
for category in le.classes_:
    px.histogram(
        top20_ingredients_words_by_category[category], y='ngram', x='count',
        orientation='h', width=600, height=400,
```

```
labels={'ngram': 'Word'}, title='{ } Top Words - Ingredients'.  
↪format(category)  
) .show()
```


cakes Top Words - Ingredients




candy Top Words - Ingredients




chocolate Top Words - Ingredients




cookies Top Words - Ingredients



seeds Top Words - Ingredients



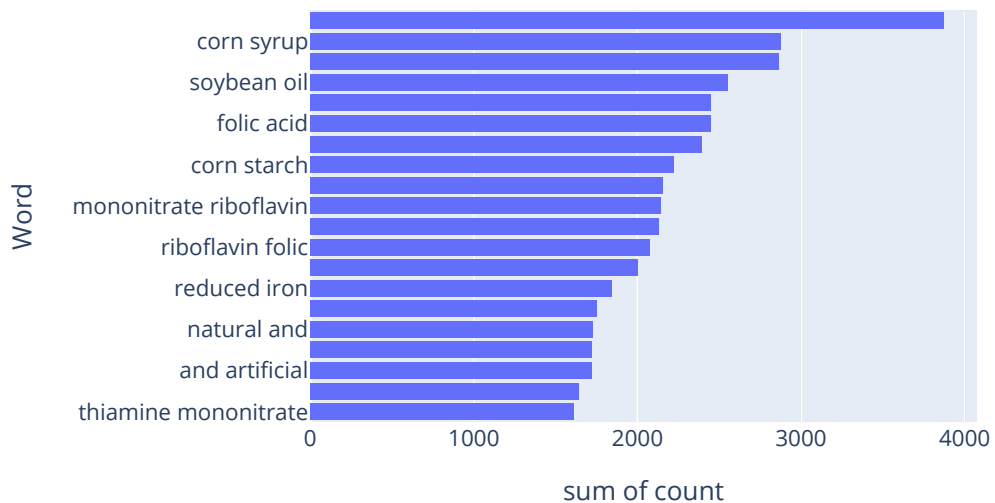
snacks Top Words - Ingredients




We can see we have a few words like 'and', 'or' in the top. Let us look at bigrams:

```
[ ]: # Do same for bigrams
eda_df, top20_ingredients_bigrams_by_category = get_ngrams_top_k(eda_df,
    ingredients', 2, 20)
for category in le.classes_:
    px.histogram(
        top20_ingredients_bigrams_by_category[category], y='ngram', x='count',
        orientation='h', width=600, height=400,
        labels={'ngram': 'Word'}, title='{} Top Bi-grams - Ingredients'.
        format(category)
    ).show()
```


cakes Top Bi-grams - Ingredients




candy Top Bi-grams - Ingredients




chocolate Top Bi-grams - Ingredients



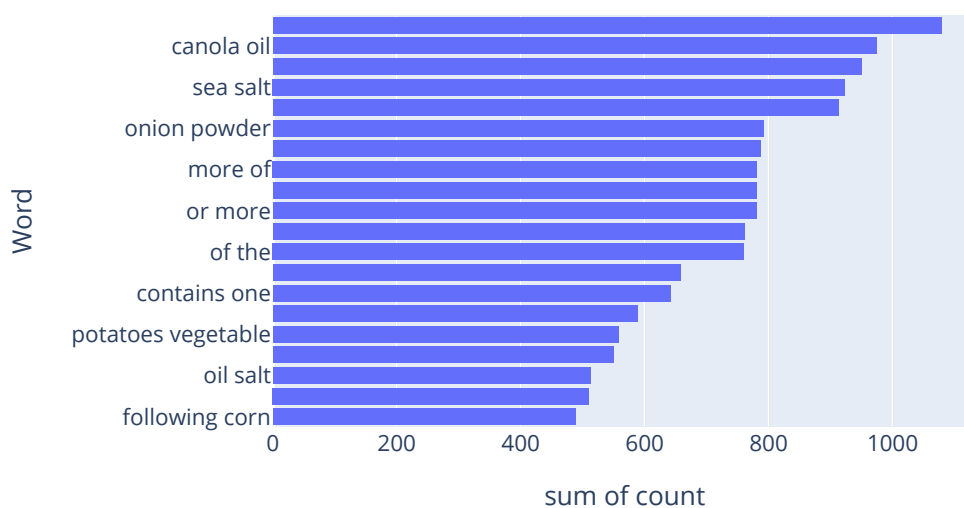
cookies Top Bi-grams - Ingredients



seeds Top Bi-grams - Ingredients




snacks Top Bi-grams - Ingredients



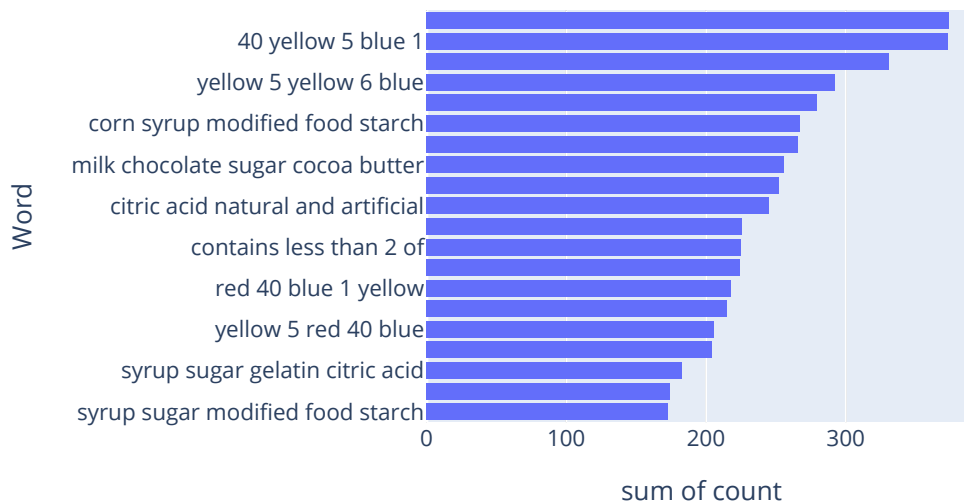
Seems we have some noise from the artificial colors and sentences. We'll keep those for now. Let us look for sentences by looking for common 5-grams:

```
[ ]: eda_df, top20_ingredients_bigrams_by_category = get_ngrams_top_k(eda_df,
    ↳ 'ingredients', 5, 20)
for category in le.classes_:
    px.histogram(
        top20_ingredients_bigrams_by_category[category], y='ngram', x='count',
    ↳ orientation='h', width=600, height=400,
        labels={'ngram': 'Word'}, title='{} Top Penta-grams - Ingredients'.
    ↳ format(category)
    ).show()
```

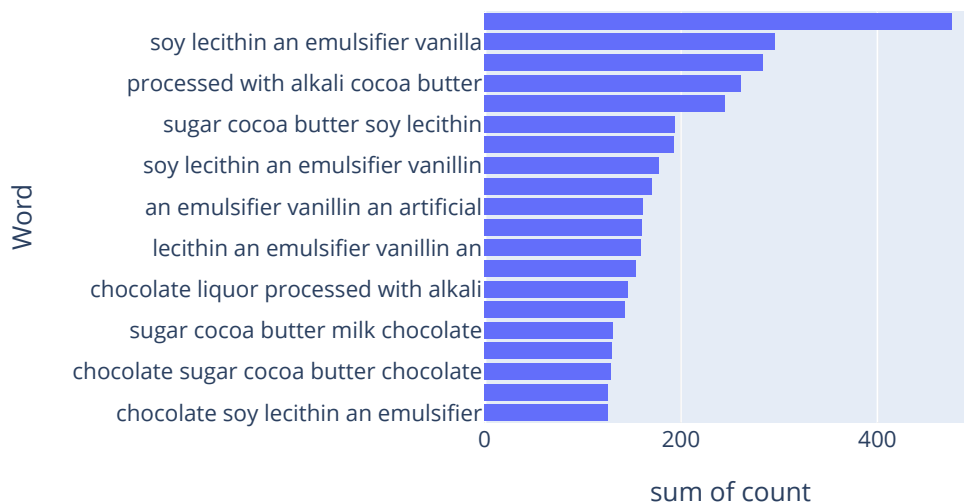
cakes Top Penta-grams - Ingredients




candy Top Penta-grams - Ingredients



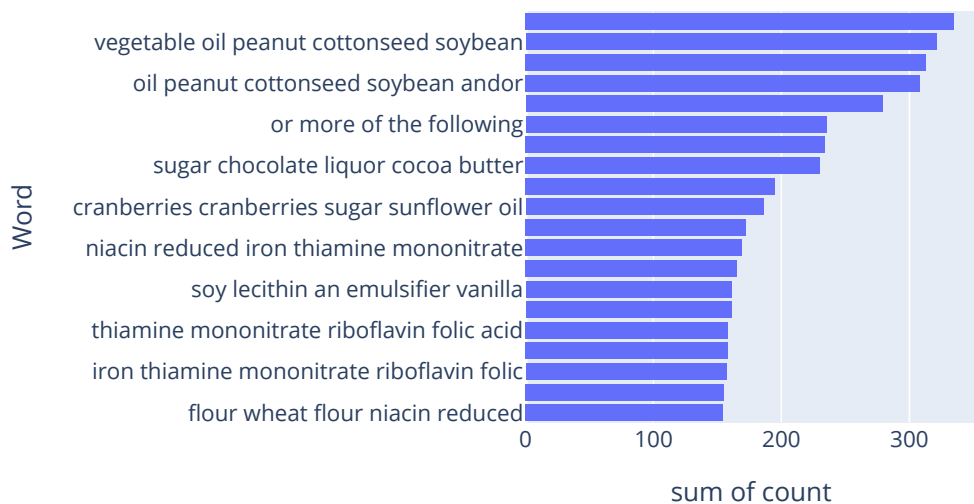
chocolate Top Penta-grams - Ingredients




cookies Top Penta-grams - Ingredients



seeds Top Penta-grams - Ingredients



snacks Top Penta-grams - Ingredients



We see that we have some sentences and numbers in the ingredients column.

```
[ ]: # eda_df, pentagrams = get_ngrams(eda_df, 'ingredients', 5)
# top_ingredient_pentagrams = sorted(pentagrams, key=pentagrams.get,
↳reverse=True)[:20]
# top_20_ingredient_pentagrams = [(key, pentagrams[key]) for key in pentagrams.
↳keys() if key in top_ingredient_pentagrams]
# top_20_ingredient_pentagrams
```

Let us do the same for the description column:


```
[ ]: eda_df, top20_description_words_by_category = get_ngrams_top_k(eda_df,
↳'description', 1, 20)
eda_df, top20_description_bigrams_by_category = get_ngrams_top_k(eda_df,
↳'description', 2, 20)
eda_df, top20_description_trigrams_by_category = get_ngrams_top_k(eda_df,
↳'description', 3, 20)
```

Top words:


```
[ ]: for category in le.classes_:
    px.histogram(
        top20_description_words_by_category[category], y='ngram', x='count',
↳orientation='h', width=600, height=400,
```

```
labels={'ngram': 'Word'}, title='{ } Top Words - Description'.  
format(category)  
) .show()
```


cakes Top Words - Description




candy Top Words - Description




chocolate Top Words - Description




cookies Top Words - Description



seeds Top Words - Description




snacks Top Words - Description




Top Bi-grams:

```
[ ]: for category in le.classes_:
    px.histogram(
        top20_description_bigrams_by_category[category], y='ngram', x='count',
        orientation='h', width=600, height=400,
        labels={'ngram': 'Word'}, title='{} Top Bi-Grams - Description'.
        format(category)
    ).show()
```


cakes Top Bi-Grams - Description




candy Top Bi-Grams - Description




chocolate Top Bi-Grams - Description




cookies Top Bi-Grams - Description



seeds Top Bi-Grams - Description



snacks Top Bi-Grams - Description



Now, let us move on to some analyses for the nutrients.


Let us first take the top 10 most common nutrients:

```
[ ]: top_10_nutrients = nutrients_column_values['name'].head(10)['value']
```

Some boxen plots by category:

```
[ ]: long = pd.melt(eda_df, id_vars=['idx', 'category', 'category_enc'],  
    ↪var_name='column')  
long = long[long['column'].isin(top_10_nutrients)].reset_index()  
boxen_plots = sns.catplot(data=long, kind='boxen', x='value', y='category',  
    ↪showfliers = False, col='column', col_wrap=5, height=3, sharex=False,  
    ↪aspect=1,  
    orient='h', margin_titles=True)  
boxen_plots.fig.subplots_adjust(top=0.85)  
boxen_plots.fig.suptitle('Boxen plots of top 10 nutrients (by frequency)'  
    ↪divided by category', fontsize=16)
```

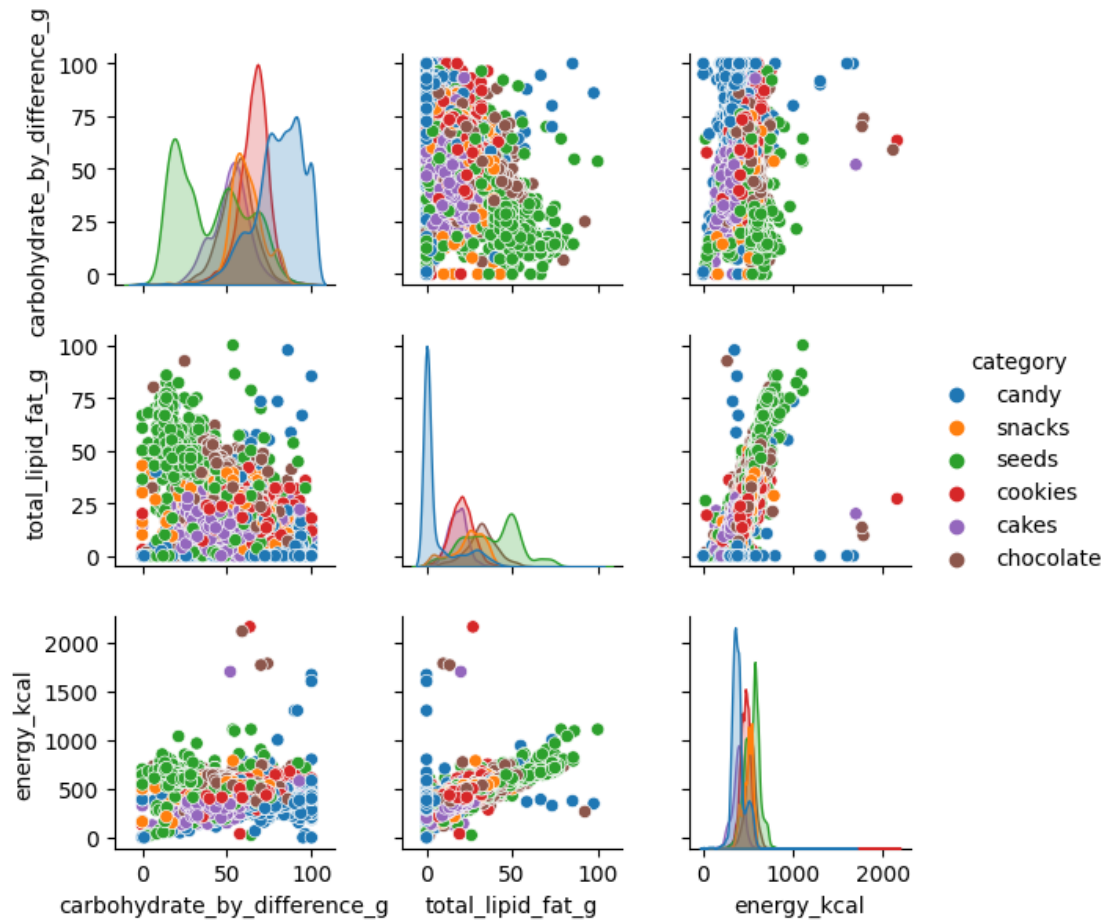
```
[ ]: Text(0.5, 0.98, 'Boxen plots of top 10 nutrients (by frequency) divided by  
category')
```



Pairwise plots by category for top 3 nutrients (to avoid a larger than necessary plot):

```
[ ]: eda_df_nutrients = eda_df[list(top_10_nutrients.head(3)) + ['category']]  
sns.pairplot(eda_df_nutrients, dropna=True, height=2, aspect=1, hue='category')
```

```
[ ]: <seaborn.axisgrid.PairGrid at 0x164ad56e940>
```



We will opt to use a neural network for the tabular data, and then use its last layer to be added to the CNN which will classify the images to get a classifier which uses both types of data.

Let us add functions which would assist us in inputting the data to a model

```
[ ]: def get_oh_categories(df, columns, sort=True):
    def sort_or_not(X):
        if not sort:
            return X
        else:
            return sorted(X)

    full_list=[]
    for column in columns:
        full_list.append(sort_or_not(df[column].unique()))
    return full_list

def get_oh_encoder(df):
```

```

cat_columns = list(get_column_values(df).keys())
categories = get_oh_categories(df, cat_columns)
ohe = preprocessing.OneHotEncoder(categories=categories, sparse=False,
↳handle_unknown='ignore')
return ohe, cat_columns

def get_oh_matrix(df):
    ohe, cat_columns = get_oh_encoder(df)
    if cat_columns:
        enc = ohe.fit_transform(df[cat_columns])
        numerical = df.select_dtypes(include=[np.number])
        result = np.concatenate([enc, numerical], axis=1)
        return result
    return df

def make_param_grids(steps, param_grids):
    final_params=[]
    # Itertools.product will do a permutation such that
    # (pca OR svd) AND (svm OR rf) will become ->
    # (pca, svm) , (pca, rf) , (svd, svm) , (svd, rf)
    for estimator_names in itertools.product(*steps.values()):
        current_grid = {}
        # Step_name and estimator_name should correspond
        # i.e preprocessor must be from pca and select.
        for step_name, estimator_name in zip(steps.keys(), estimator_names):
            for param, value in param_grids.get(estimator_name).items():
                if param == 'object':
                    # Set actual estimator in pipeline
                    current_grid[step_name]=[value]
                else:
                    # Set parameters corresponding to above estimator
                    current_grid[step_name+'__'+param]=value
            #Append this dictionary to final params
            final_params.append(current_grid)
    return final_params

def create_features_from_strings(df, column, strings):
    if type(strings) == str:
        strings = [strings]
    for string in strings:
        mask = df[column].str.contains(string)
        new_column = (mask * 1).to_frame().rename(columns={column: '{}_{}'.format(column, string.replace(' ', '_'))}) # Doing it like this because
↳apparently

```

```

        df = pd.concat([df, new_column], axis=1)
        # otherwise performance is hurt
    ↪
    return df

def fix_brands(df):
    brands_regex_dict = {
        'walmart': ['walmart'],
        'target': ['target'],
        'ferrara': ['ferrara'],
        'whole foods': ['whole foods'],
        'safeway': ['safeway'],
        'walgreens': ['walgreens'],
        'mars chocolate': ['mars chocolate'],
        'lindt': ['lindt'],
        'star snacks': ['star snacks'],
        'north star': ['north star'],
        'weis': ['weis'],
        'harristeeter': ['harristeeter', 'harris teeter'],
        'other': ['not a branded item']
    }
    for string in brands_regex_dict:
        mask = df['brand'].str.contains(''.join(brands_regex_dict[string]))
        df.loc[mask, 'brand'] = string
    return df

def fe(df, feature_strings_dict, drop_brand=True):
    for column, dict in feature_strings_dict.items():
        df = create_features_from_strings(df, column, dict)
    df = df.drop(['description', 'ingredients', 'serving_size_unit',
    ↪ 'household_serving_fulltext'], axis=1)
    if drop_brand:
        df = df.drop(['brand'], axis=1)
    return df

```

First of all, let us remove all columns with more than 80% nulls

```

[ ]: nulls = X.isnull().sum().to_frame().reset_index().rename(columns={'index':
    ↪ 'column', 0: 'nulls'})
    nulls['pc'] = nulls['nulls'] / len(X) * 100
    columns_to_drop = nulls[nulls['pc'] >= 80]['column'].to_list()

```

So we have 25.4K records in our training set;

0.05 of them to find the best way to fill missing data

0.3 for feature engineering, type of model, and 0.35 model tuning

```
[ ]: columns_to_drop = columns_to_drop + ['category', 'category_enc']
X = clean_text_values(dataset_w_nutrients.drop(columns_to_drop, axis=1)).
    ↪set_index('idx')
y = dataset_w_nutrients[['idx', 'category_enc']].
    ↪set_index('idx')['category_enc']
X_train, X_test, y_train, y_test = model_selection.train_test_split(X, y,
    ↪test_size=0.2, random_state=1337, stratify=y)

# 0.05 for missing data imputation
X_rest, X_rest2, y_rest, y_rest2 = model_selection.train_test_split(X_train,
    ↪y_train, test_size=0.1, random_state=1337, stratify=y_train)
X_mdi, X_coc, y_mdi, y_coc = model_selection.train_test_split(X_rest2, y_rest2,
    ↪test_size=0.5, random_state=1337, stratify=y_rest2)

# 0.3 for feature engineering, resampling, and 0.35 model tuning
X_rest, X_fe, y_rest, y_fe = model_selection.train_test_split(X_rest, y_rest,
    ↪test_size=1/3, random_state=1337, stratify=y_rest)
X_rs, X_mt, y_rs, y_mt = model_selection.train_test_split(X_rest, y_rest,
    ↪test_size=0.5, random_state=1337, stratify=y_rest)
X_mt, y_mt = pd.concat([X_mt, X_coc], axis=0), pd.concat([y_mt, y_coc], axis=0)
```

Initial feature engineering

First, let's clean the text columns.

We saw in the EDA that there are some sentences in the ingredients column, as well as numbers.


Let us look for those sentences again:

Let us extract a list of most common ingredients from each category to look for fishy strings and get some ideas for features:


```
[ ]: X_fe_with_cat = pd.concat([X_fe.reset_index(), pd.Series(1e.
    ↪inverse_transform(y_fe))], axis=1).rename(columns={0: 'category'})

[ ]: X_fe_with_cat, top_15_words = get_ngrams_top_k(X_fe_with_cat, 'ingredients', 1,
    ↪15, True)
for cat, top in top_15_words.items():
    px.histogram(top, orientation='h', y='ngram', x='count', title='{ } top_
    ↪words'.format(cat), height=400, width=500).show()
```


seeds top words




chocolate top words




candy top words




snacks top words



cookies top words



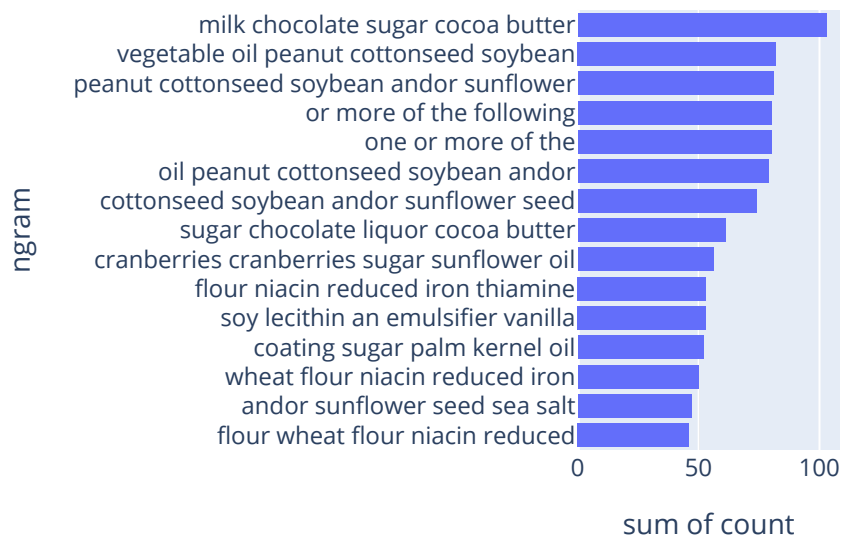
cakes top words



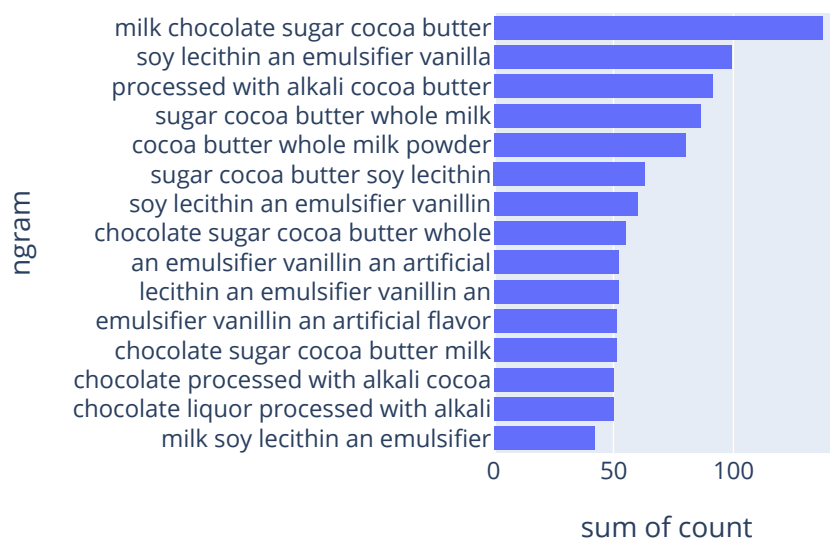
Let's do the same for 5-grams to identify common sentences:

```
[ ]: X_fe_with_cat, top_15_5grams = get_ngrams_top_k(X_fe_with_cat, 'ingredients', 5, 15, True)
for cat, top in top_15_5grams.items():
    px.histogram(top, orientation='h', y='ngram', x='count', title='{cat} top 5-grams'.format(cat), height=400, width=500).show()
```


seeds top 5-grams




chocolate top 5-grams




candy top 5-grams




snacks top 5-grams



cookies top 5-grams



cakes top 5-grams




We can see that we have some bad words like 'or', 'and' etc (also numbers). Let's fix that:

```
[ ]: irrelevant_strings = ['and', 'contains', 'one', 'or', 'more', 'less', 'than',
    ↪ 'of', 'a', 'the', 'following'] + [i for i in range(0, 1000)]
for word in irrelevant_strings:
    X_fe['ingredients'] = X_fe['ingredients'].str.replace("{} {}".format(word),
    ↪ " ")
    X_fe_with_cat['ingredients'] = X_fe_with_cat['ingredients'].str.replace("{}
    ↪ {}".format(word), " ")
```


Let's see the new top words and 5-grams:

```
[ ]: X_fe_with_cat, top_15_words = get_ngrams_top_k(X_fe_with_cat, 'ingredients', 1,
    ↪ 15, True)
for cat, top in top_15_words.items():
    px.histogram(top, orientation='h', y='ngram', x='count', title='{} top
    ↪ words'.format(cat), height=400, width=500).show()
```


seeds top words




chocolate top words




candy top words




snacks top words



cookies top words

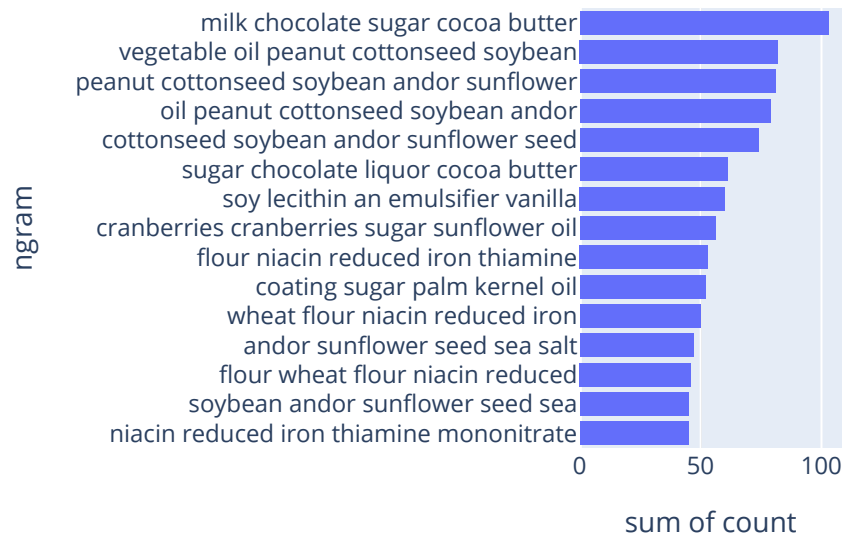


cakes top words

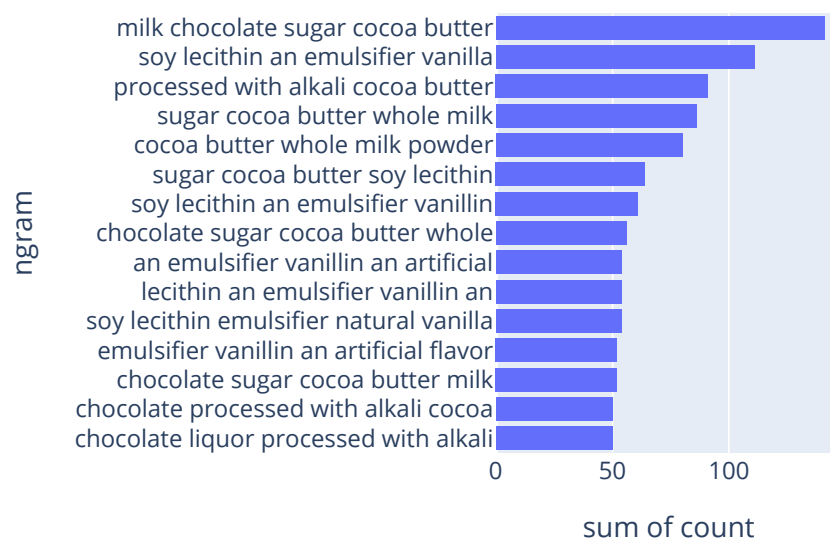


```
[ ]: X_fe_with_cat, top_15_5grams = get_ngrams_top_k(X_fe_with_cat, 'ingredients', 5, 15, True)
for cat, top in top_15_5grams.items():
    px.histogram(top, orientation='h', y='ngram', x='count', title='{} top_5-grams'.format(cat), height=400, width=500).show()
```


seeds top 5-grams




chocolate top 5-grams




candy top 5-grams




snacks top 5-grams



cookies top 5-grams



cakes top 5-grams



Numbers and sentences eliminated!


Let's move on to create features from the ingredients:

```
[ ]: ingredients_strings = np.unique(np.concatenate([top_15_words[cat]['ngram'].
    ↪ values for cat in top_15_words.keys()], axis=0))
```


Let's look for common words in the description

```
[ ]: X_fe_with_cat, top_15_desc_words = get_ngrams_top_k(X_fe_with_cat,
    ↪ 'description', 1, 15)
for cat, top in top_15_desc_words.items():
    px.histogram(top, orientation='h', y='ngram', x='count', title='{ } top
    ↪ words'.format(cat), height=400, width=500).show()
```


seeds top words




chocolate top words




candy top words




snacks top words



cookies top words



cakes top words




This time we'll compile the strings manually because we don't need all of them:

```
[ ]: description_strings = np.array([
    'roasted', 'mix', 'almonds', 'corn', 'nuts', 'chocolate', 'trail',
    'cashews', 'salt', 'covered', 'dark', 'milk', 'caramel', 'bar',
    'truffles', 'candy', 'fruit', 'sour', 'gumm', 'jell', 'candies', 'cherry',
    'chewy', 'chips', 'potato', 'tortilla', 'kettle', 'pretzel',
    'cookie', 'chip', 'sugar', 'sandwich', 'butter', 'oatmeal', 'vanilla',
    'fudge', 'cake', 'pie', 'donut', 'cupcakes', 'cheesecake', 'bakery',
    'brownie'
])
```


Let's do the same for the brands:

```
[ ]: X_fe_with_cat, top_15_brand_words = get_ngrams_top_k(X_fe_with_cat, 'brand', 1,
    15)
for cat, top in top_15_brand_words.items():
    px.histogram(top, orientation='h', y='ngram', x='count', title='{ } top
    words'.format(cat), height=400, width=500).show()
```


seeds top words




chocolate top words




candy top words




snacks top words



cookies top words



cakes top words




```
[ ]: brand_strings = np.array([
    'snack', 'chocola', 'candy', 'candies', 'chips', 'potato', 'bake',
    'baking', 'biscuit', 'cookie', 'jelly', 'nuts'
])
```


And do the same for the household_serving_fulltext:

```
[ ]: X_fe_with_cat, top_15_household_serving_fulltext_words =
    get_ngrams_top_k(X_fe_with_cat, 'household_serving_fulltext', 1, 15)
for cat, top in top_15_household_serving_fulltext_words.items():
    px.histogram(top, orientation='h', y='ngram', x='count', title='{cat} top
    words'.format(cat), height=400, width=500).show()
```


seeds top words




chocolate top words




candy top words




snacks top words



cookies top words



cakes top words




```
[ ]: household_serving_fulltext_strings = np.array([
    'cup', 'onz', 'tbsp', 'pieces', 'about', 'package', 'grm', 'bag', 'bar',
    'piece', 'squares', 'pcs', 'pouch', 'chips', 'pretzels', 'cookie',
    'wafers', 'crackers', 'slice', 'pie', 'cake', 'cupcake', 'donut'
])
```


Let us do the same for bi-grams:

```
[ ]: X_fe_with_cat, top_7_ingredients_bigrams = get_ngrams_top_k(X_fe_with_cat,
    'ingredients', 2, 7, True)
for cat, top in top_7_ingredients_bigrams.items():
    px.histogram(top, orientation='h', y='ngram', x='count', title='{ } top
    words'.format(cat), height=400, width=500).show()
    ingredients_strings = np.concatenate([ingredients_strings, top['ngram'].
    values], axis=0)
```


seeds top words




chocolate top words




candy top words




snacks top words



cookies top words




cakes top words




```
[ ]: X_fe_with_cat, top_7_description_bigrams = get_ngrams_top_k(X_fe_with_cat,
    ↳ 'description', 2, 7, True)
for cat, top in top_7_description_bigrams.items():
    px.histogram(top, orientation='h', y='ngram', x='count', title='{ } top_
    ↳ words'.format(cat), height=400, width=500).show()
    description_strings = np.concatenate([description_strings, top['ngram'].
    ↳ values], axis=0)
```


seeds top words




chocolate top words




candy top words




snacks top words



cookies top words



cakes top words



Let us create a proper dictionary for convinience later:

```
[ ]: feature_strings_dict = {  
    'ingredients': np.unique(ingredients_strings),  
    'description': np.unique(description_strings),  
    'brand': np.unique(brand_strings),  
    'household_serving_fulltext': np.unique(household_serving_fulltext_strings)  
}
```

We also created a function to fix some brand values (which were manually picked from the X_fe dataset).

After the initial feature engineering, let us choose how we will impute our values

We'll create a pipeline which first imputes the missing values with various methods, and then cross validates the results of a vanilla XGB classifier. The methods used are comprised of simple imputations including the mean, median values or filling all missing values with 0, and more complex imputations using linear and trees regressions with various parameters to impute the missing values. Using balanced accuracy score to account for class imbalance, as we have yet to address it.

```
[ ]: X_mdi_fe = fe(X_mdi, feature_strings_dict)
```

```
[ ]: mdi_pipeline_steps = {'imputer':['simple_imp', 'iterative_imp'],
                           'classifier':['xgbc']}

# fill parameters to be searched in this dict
mdi_all_param_grids = {'xgbc':{'object':XGBClassifier(random_state=1337),
                              'use_label_encoder': [False],
                              'objective': ['multi:softmax'],
                              'eval_metric': ['mlogloss'],
                              , 'tree_method': ['gpu_hist']
                              },
                      'simple_imp':{'object': SimpleImputer(),
                                    'strategy':['mean', 'median', 'constant']}
                      , 'iterative_imp':{'object':IterativeImputer(random_state=1337),
                                          'estimator':[linear_model.
↳ElasticNet(random_state=1337)],
                                          'max_iter': [200],
                                          'skip_complete': [True],
                                          'initial_strategy': ['mean', 'median', 'constant'],
                                          'imputation_order': ['ascending', 'descending'],
↳'random']
                      }

mdi_param_grids_list = make_param_grids(mdi_pipeline_steps, mdi_all_param_grids)
mdi_pipe = Pipeline(steps=[('imputer', SimpleImputer()), ('classifier',
↳XGBClassifier())])
mdi_grid = model_selection.GridSearchCV(mdi_pipe, param_grid =
↳mdi_param_grids_list, scoring='balanced_accuracy', cv=10, verbose=15)
mdi_grid.fit(X_mdi_fe, y_mdi)
```

```
[ ]: mdi_grid.best_params_
```

```
[ ]: pd.DataFrame(mdi_grid.cv_results_).sort_values('rank_test_score').
↳head(5)[['params', 'mean_test_score']].values
```

Result is:

```
IterativeImputer(estimator=ElasticNet(random_state=1337), imputation_order='descending',
initial_strategy='median', max_iter=200, random_state=1337, skip_complete=True)
```

On to resampling.

We won't test only downsampling because we want to get a balanced dataset, and the minority classes have too few samples.

We'll test several strategies - Random upsampling, random upsampling with downsampling,

SMOTE upsampling, SMOTE upsampling with downsampling, KMeans SMOTE oversampling, KMeans SMOTE oversampling with downsampling.

For downsampling we will test random downsampling, and no downsampling.

Each strategy which involves downsampling will test 10 values - 0.95, 0.9, 0.85, 0.8, 0.75, 0.7, 0.65, 0.6, 0.55, 0.5 of the largest class

The random upsampling strategies will also test shrinkage values - 0, 1, 2, 3

The SMOTE strategies will test different numbers of k_neighbors - 5, 10, 15

In total we get 77 strategies to choose from.

In this case we will also use a vanilla XGBoost to compare out results (this time using accuracy instead of balanced accuracy)

```
[ ]: chosen_imputer = IterativeImputer(estimator=linear_model.  
    ↪ElasticNet(random_state=1337),  
        imputation_order='descending', initial_strategy='mean',  
        max_iter=200, random_state=1337, skip_complete=True)  
chosen_imputer.fit(X_mdi_fe)
```

```
[ ]: IterativeImputer(estimator=ElasticNet(random_state=1337),  
    imputation_order='descending', max_iter=200, random_state=1337,  
    skip_complete=True)
```

```
[ ]: X_rs_fe = fe(X_rs, feature_strings_dict)  
X_rs_imputed = chosen_imputer.transform(X_rs_fe)
```

```
[ ]: # Had to give up lambdas to recognize functions after CV :(  
  
# # Creates lambdas to feed into the strategies argument of the downsamplers ↵  
↪(because of the CV the numbers are not constant)  
# def get_strategy_lambda(i):  
#     return lambda y: {j: round(max(np.bincount(y)) * i * 0.1) for j in ↵  
↪range(0,6)}  
# strategy_lambdas = [get_strategy_lambda(i) for i in range(6, 10)]  
strategy_funcs = []  
def_strategy_function_exec_code = """def strategy_{}(y): return {{j: ↵  
    ↪round(max(np.bincount(y)) * {} * 0.05) for j in range(0,6)}}"""  
for i in range(10,20):  
    exec(def_strategy_function_exec_code.format(round(i * 5), i))  
    exec("""strategy_funcs.append(strategy_{})""".format(i * 5))  
  
# Had to use these step names becasue imblearn uses a custom pipe constructor.  
# found this out after I finished these dicts and almost cried in fear of a ↵  
↪much more complicated CV  
rs_pipeline_steps = {
```

```

        'randomoversampler': ['random_us', 'smote_us'],
        'randomundersampler': ['no_ds', 'random_ds'],
        'xgbclassifier': ['xgbc']
    }

rs_all_param_grids = {'xgbc': {'object': XGBClassifier(random_state=1337),
                              'use_label_encoder': [False],
                              'objective': ['multi:softmax'],
                              'eval_metric': ['mlogloss'],
                              , 'tree_method': ['gpu_hist']
                              },
                      'random_us': {
                          'object': over_sampling.
↳ RandomOverSampler(random_state=1337),
                          'shrinkage': [0, 1, 2, 3]
                      },
                      'smote_us': {
                          'object': over_sampling.SMOTE(random_state=1337),
                          'k_neighbors': [5, 10, 15]
                      },
                      'no_ds': {
                          'object': None
                      },
                      'random_ds': {
                          'object': under_sampling.
↳ RandomUnderSampler(random_state=1337),
                          'sampling_strategy': strategy_funcs
                      }
    }

rs_param_grids_list = make_param_grids(rs_pipeline_steps, rs_all_param_grids)
# rs_pipe = Pipeline(steps=[('up_sample', None), ('down_sample', None),
↳ ('classifier', None)])
rs_pipe = pipeline.make_pipeline(over_sampling.RandomOverSampler(),
↳ under_sampling.RandomUnderSampler(), XGBClassifier()) # imblearn pipe
rs_grid = model_selection.GridSearchCV(rs_pipe, param_grid =
↳ rs_param_grids_list, scoring='accuracy', cv=10, verbose=15)
rs_grid.fit(X_rs_imputed, y_rs)

```

Results are: SMOTE(k_neighbors=10, random_state=1337)

RandomUnderSampler(random_state=1337, sampling_strategy=<function strategy_75 at 0x7fec5bb6def0>)

So, SMOTE with k_neighbors=10 and RandomUnderSampler with a downsampling to 75% of the largest class.

```
[ ]: pd.options.mode.chained_assignment = None

brands_for_column_values = X_fe[['brand']] # Take brands from FE portion of
↳dataset
brands_column_values = get_column_values(brands_for_column_values)
brands = X_fe[['brand']]
ok_brands = brands_column_values['brand'][brands_column_values['brand']['pc']_
↳>= 0.1]['value'].values
brands_mask = ~brands['brand'].isin(ok_brands)
brands.loc[brands_mask, 'brand'] = 'other'
brands_ohe, _ = get_oh_encoder(brands)
brands_ohe.fit(brands)
```

```
[ ]: OneHotEncoder(categories=[['7eleven inc', 'abimar foods inc', 'ahold usa inc',
                                'aldibenner company', 'american halal company inc',
                                'american importing co inc', 'archer farms',
                                'back to nature foods company llc',
                                'bee international inc', 'bergin nut company inc',
                                'best choice', 'better made snack foods inc',
                                'big y foods inc', 'bimbo bakeries usa inc',
                                'bjs wholesale club corporate brands',
                                'blue diamond growers', 'brads raw chips',
                                'brookshire grocery company',
                                'california flavored nuts', 'candyrific llc',
                                'cape cod potato chips inc', 'charms company',
                                'cibo vita inc', 'creative natural products inc',
                                'creative snacks co llc',
                                'csm bakery products na inc', 'cvs pharmacy inc',
                                'dawn food products inc', 'delhaize america inc',
                                'demets candy company', ...]],
                    handle_unknown='ignore', sparse=False)
```

Let us remove highly correlated features

Now let us write a function which ties the dataset together, before moving to high correlated columns removal

```
[ ]: def get_numerical_df(df, feature_strings_dict, brands_ohe):
    df = fix_brands(df)
    brands = df[['brand']]
    brands_oh = brands_ohe.transform(brands)
    df_fe = fe(df, feature_strings_dict)
    numerical_df = np.concatenate([brands_oh, df_fe], axis=1)
    return numerical_df
```

```
[ ]: # Refit chosen imputer to the new columns
X_mdi_numerical = get_numerical_df(X_mdi, feature_strings_dict, brands_ohe)
```

```

chosen_imputer = IterativeImputer(estimator=linear_model.
    ↪ElasticNet(random_state=1337),
                                imputation_order='descending', initial_strategy='mean',
                                max_iter=200, random_state=1337, skip_complete=True)
chosen_imputer.fit(X_mdi_numerical)

```

```

[ ]: IterativeImputer(estimator=ElasticNet(random_state=1337),
                        imputation_order='descending', max_iter=200, random_state=1337,
                        skip_complete=True)

```

```

[ ]: X_fe_numerical = get_numerical_df(X_fe, feature_strings_dict, brands_ohe)
X_fe_numerical_imputed = chosen_imputer.transform(X_fe_numerical)

```

Let's use CV to find best value to consider a high correlation:

```

[ ]: def strategy_75(y):
    return {j: round(max(np.bincount(y)) * 15 * 0.05) for j in range(0,6)}

chosen_os = over_sampling.SMOTE(k_neighbors=10, random_state=1337)
chosen_us = under_sampling.RandomUnderSampler(random_state=1337,
    ↪sampling_strategy=strategy_75)
vanilla_xgbc = XGBClassifier(eval_metric='mlogloss', objective='multi:softmax',
                             random_state=1337, tree_method='gpu_hist',
                             use_label_encoder=False)

```

```

[ ]: corr_pipe = pipeline.make_pipeline(chosen_os, chosen_us, vanilla_xgbc)
cv_scores = {}
for corr in [0.6, 0.7, 0.8, 0.9]:
    X_fe_numerical_df = pd.DataFrame(X_fe_numerical_imputed)
    X_fe_corr = X_fe_numerical_df.corr()
    upper_tri = X_fe_corr.where(np.triu(np.ones(X_fe_corr.shape), k=1).
    ↪astype(bool))
    corr_columns_to_drop = [column for column in upper_tri.columns if
    ↪any(upper_tri[column] > corr)]
    X_fe_numerical_corr = X_fe_numerical_df.drop(corr_columns_to_drop, axis=1).
    ↪values
    cv_scores[corr] = model_selection.cross_val_score(corr_pipe,
    ↪X_fe_numerical_corr, y_fe, cv=10,
                                                    verbose=15, n_jobs=-1).
    ↪mean()
max_score = max(cv_scores, key=cv_scores.get)
max_score

```

So we will remove all columns with correlation above the chosen threshold.

```

[ ]: X_fe_numerical_df = pd.DataFrame(X_fe_numerical_imputed)
X_fe_corr = X_fe_numerical_df.corr()

```

```
upper_tri = X_fe_corr.where(np.triu(np.ones(X_fe_corr.shape), k=1).astype(bool))
corr_columns_to_drop = [column for column in upper_tri.columns if
    ↪any(upper_tri[column] > max_score)]
```

```
[ ]: def get_final_matrix(X, feature_strings_dict, brands_ohe, corr_columns_to_drop,
    ↪chosen_imputer):
    numerical_X = get_numerical_df(X, feature_strings_dict, brands_ohe)
    X_imputed_df = pd.DataFrame(chosen_imputer.transform(numerical_X))
    final_matrix = X_imputed_df.drop(corr_columns_to_drop, axis=1).values
    return final_matrix
```

```
[ ]: X_mt_numerical = get_final_matrix(X_mt, feature_strings_dict, brands_ohe,
    ↪corr_columns_to_drop, chosen_imputer)
```

Now let us tune our XGB Classifier one step at a time, starting with the number of trees for the default learning rate.

```
[ ]: mt_pipeline_steps_ne = {
    'xgbclassifier': ['xgbc']
}

mt_all_param_grids_ne = {'xgbc': {'object': XGBClassifier(random_state=1337),
    'n_estimators': [100 + i * 10 for i in range(0, 41)],
    'use_label_encoder': [False],
    'objective': ['multi:softmax'],
    'eval_metric': ['mlogloss'],
    'tree_method': ['gpu_hist']
    },
}

mt_param_grids_list_ne = make_param_grids(mt_pipeline_steps_ne,
    ↪mt_all_param_grids_ne)
mt_pipe_ne = pipeline.make_pipeline(chosen_os, chosen_us, vanilla_xgbc)
mt_grid_ne = model_selection.GridSearchCV(mt_pipe_ne, param_grid =
    ↪mt_param_grids_list_ne, scoring='accuracy', cv=10, verbose=15)
mt_grid_ne.fit(X_mt_numerical, y_mt)
```

```
[ ]: mt_grid_ne.best_params_['xgbclassifier__n_estimators']
```

Now let's tune other parameters:

```
[ ]: mt_pipeline_steps_mdcw = {
    'xgbclassifier': ['xgbc']
}

mt_all_param_grids_mdcw = {'xgbc': {'object': XGBClassifier(random_state=1337),
```



```

        'n_estimators': [mt_grid_ne.
↪best_params_['xgbclassifier__n_estimators']],
        'max_depth': [i for i in range(3, 13)],
        'min_child_weight': [i for i in range(1, 7)],
        'use_label_encoder': [False],
        'objective': ['multi:softmax'],
        'eval_metric': ['mlogloss']
    , 'tree_method': ['gpu_hist']
    }

    }

mt_param_grids_list_mdcw = make_param_grids(mt_pipeline_steps_mdcw,
↪mt_all_param_grids_mdcw)
mt_pipe_mdcw = pipeline.make_pipeline(chosen_os, chosen_us, vanilla_xgbc)
mt_grid_mdcw = model_selection.GridSearchCV(mt_pipe_mdcw, param_grid =
↪mt_param_grids_list_mdcw, scoring='accuracy', cv=10, verbose=15)
mt_grid_mdcw.fit(X_mt_numerical, y_mt)

```

```
[ ]: mt_grid_mdcw.best_params_
```

```
[ ]: print(mt_grid_mdcw.best_params_['xgbclassifier__max_depth'])
print(mt_grid_mdcw.best_params_['xgbclassifier__min_child_weight'])
```

```
[ ]: mt_pipeline_steps_gam = {
    'xgbclassifier': ['xgbc']
}

mt_all_param_grids_gam = {'xgbc': {'object': XGBClassifier(random_state=1337),
    'n_estimators': [mt_grid_ne.
↪best_params_['xgbclassifier__n_estimators']],
    'max_depth': [mt_grid_mdcw.
↪best_params_['xgbclassifier__max_depth']],
    'min_child_weight': [mt_grid_mdcw.
↪best_params_['xgbclassifier__min_child_weight']],
    'gamma': [i / 10.0 for i in range(0, 5)],
    'use_label_encoder': [False],
    'objective': ['multi:softmax'],
    'eval_metric': ['mlogloss']
    , 'tree_method': ['gpu_hist']
    }

    }

mt_param_grids_list_gam = make_param_grids(mt_pipeline_steps_gam,
↪mt_all_param_grids_gam)
mt_pipe_gam = pipeline.make_pipeline(chosen_os, chosen_us, vanilla_xgbc)

```

```
mt_grid_gam = model_selection.GridSearchCV(mt_pipe_gam, param_grid =  
    ↪mt_param_grids_list_gam, scoring='accuracy', cv=10, verbose=15)  
mt_grid_gam.fit(X_mt_numerical, y_mt)
```

```
[ ]: mt_grid_gam.best_params_['xgbclassifier__gamma']
```

On to check subsample and colsample_bytree.

We will first check with 0.1 intervals, and then 0.05 intervals around the chosen values.

```
[ ]: mt_pipeline_steps_sscs = {  
    'xgbclassifier': ['xgbc']  
}  
  
mt_all_param_grids_sscs = {'xgbc': {'object': XGBClassifier(random_state=1337),  
    'n_estimators': [mt_grid_ne.  
    ↪best_params_['xgbclassifier__n_estimators']],  
    'max_depth': [mt_grid_mdcw.  
    ↪best_params_['xgbclassifier__max_depth']],  
    'min_child_weight': [mt_grid_mdcw.  
    ↪best_params_['xgbclassifier__min_child_weight']],  
    'gamma': [mt_grid_gam.  
    ↪best_params_['xgbclassifier__gamma']],  
    'subsample': [i/10.0 for i in range(6,11)],  
    'colsample_bytree': [i/10.0 for i in range(6,11)],  
    'use_label_encoder': [False],  
    'objective': ['multi:softmax'],  
    'eval_metric': ['mlogloss'],  
    'tree_method': ['gpu_hist']  
    },  
    }  
  
mt_param_grids_list_sscs = make_param_grids(mt_pipeline_steps_sscs,  
    ↪mt_all_param_grids_sscs)  
mt_pipe_sscs = pipeline.make_pipeline(chosen_os, chosen_us, vanilla_xgbc)  
mt_grid_sscs = model_selection.GridSearchCV(mt_pipe_sscs, param_grid =  
    ↪mt_param_grids_list_sscs, scoring='accuracy', cv=10, verbose=15)  
mt_grid_sscs.fit(X_mt_numerical, y_mt)
```

```
[ ]: ss = mt_grid_sscs.best_params_['xgbclassifier__subsample']  
csbt = mt_grid_sscs.best_params_['xgbclassifier__colsample_bytree']  
print(ss)  
print(csbt)
```

```
[ ]: ss = mt_grid_sscs.best_params_['xgbclassifier__subsample']  
csbt = mt_grid_sscs.best_params_['xgbclassifier__colsample_bytree']  
print(ss)
```

```
print(csbt)
```

```
[ ]: mt_pipeline_steps_sscs2 = {
    'xgbclassifier': ['xgbc']
}

mt_all_param_grids_sscs2 = {'xgbc': {'object': XGBClassifier(random_state=1337),
    'n_estimators': [mt_grid_ne.
↳best_params_['xgbclassifier__n_estimators']],
    'max_depth': [mt_grid_mdcw.
↳best_params_['xgbclassifier__max_depth']],
    'min_child_weight': [mt_grid_mdcw.
↳best_params_['xgbclassifier__min_child_weight']],
    'subsample': [ss - 0.05, ss, ss + 0.05],
    'colsample_bytree': [csbt - 0.05, csbt, csbt + 0.5],
    'use_label_encoder': [False],
    'objective': ['multi:softmax'],
    'eval_metric': ['mlogloss']
    , 'tree_method': ['gpu_hist']
    }

    }

mt_param_grids_list_sscs2 = make_param_grids(mt_pipeline_steps_sscs2,↳
↳mt_all_param_grids_sscs2)

mt_pipe_sscs2 = pipeline.make_pipeline(chosen_os, chosen_us, vanilla_xgbc)
mt_grid_sscs2 = model_selection.GridSearchCV(mt_pipe_sscs2, param_grid =↳
↳mt_param_grids_list_sscs2, scoring='accuracy', cv=10, verbose=15)
mt_grid_sscs2.fit(X_mt_numerical, y_mt)
```

```
[ ]: mt_grid_sscs2.best_params_['xgbclassifier__subsample']
```

Now for regularization parameters:

```
[ ]: mt_pipeline_steps_reg = {
    'xgbclassifier': ['xgbc']
}

mt_all_param_grids_reg = {'xgbc': {'object': XGBClassifier(random_state=1337),
    'n_estimators': [mt_grid_ne.
↳best_params_['xgbclassifier__n_estimators']],
    'max_depth': [mt_grid_mdcw.
↳best_params_['xgbclassifier__max_depth']],
    'min_child_weight': [mt_grid_mdcw.
↳best_params_['xgbclassifier__min_child_weight']],
    'subsample': [mt_grid_sscs2.
↳best_params_['xgbclassifier__subsample']],
```

```

        'colsample_bytree': [mt_grid_sscs2.
↪best_params_['xgbclassifier__colsample_bytree']],
        'alpha': [0, 1e-5, 1e-2, 0.1, 1],
        'lambda': [0, 1e-5, 1e-2, 0.1, 1],
        'use_label_encoder': [False],
        'objective': ['multi:softmax'],
        'eval_metric': ['mlogloss']
    , 'tree_method': ['gpu_hist']
    }

    }

mt_param_grids_list_reg = make_param_grids(mt_pipeline_steps_reg,
↪mt_all_param_grids_reg)
mt_pipe_reg = pipeline.make_pipeline(chosen_os, chosen_us, vanilla_xgbc)
mt_grid_reg = model_selection.GridSearchCV(mt_pipe_reg, param_grid =
↪mt_param_grids_list_reg, scoring='accuracy', cv=10, verbose=15)
mt_grid_reg.fit(X_mt_numerical, y_mt)

```

```
[ ]: mt_grid_reg.best_params_
```

For some reason tuning the learning rate parameter didn't work (got exact same results in CV).

So now we have out XGB classifier!

```

XGBClassifier(alpha=0, colsample_bytree=0.8, eval_metric='mlogloss', lambda=0,
max_depth=12, min_child_weight=3, n_estimators=250, objective='multi:softmax', ran-
dom_state=1337, subsample=1.0, tree_method='gpu_hist', use_label_encoder=False)

```

Let us move on to choosing a RF classifier using CV:

```

[ ]: mt_pipeline_steps_rf = {
    'randomforestclassifier': ['rf']
}

mt_all_param_grids_rf = {'rf': {'object': ensemble.
↪RandomForestClassifier(random_state=1337),
    'n_estimators': [100 + i * 50 for i in range(0, 9)],
    'criterion': ['entropy'],
    'max_features': ['sqrt'],
    'min_samples_split': [2, 3, 4, 5]
    }
    }

mt_param_grids_list_rf = make_param_grids(mt_pipeline_steps_rf,
↪mt_all_param_grids_rf)
mt_pipe_rf = pipeline.make_pipeline(chosen_os, chosen_us, vanilla_rf)
mt_grid_rf = model_selection.GridSearchCV(mt_pipe_rf, param_grid =
↪mt_param_grids_list_rf, scoring='accuracy', cv=10, verbose=15)

```

```
mt_grid_rf.fit(X_mt_numerical, y_mt)
```

So, our chosen RF classifier is is:

```
RandomForestClassifier(criterion='entropy', n_estimators=200, random_state=1337)
```

Let us move on to a SVM classifier:

```
[ ]: mt_pipeline_steps_svm = {
    'svc': ['svc']
}

mt_all_param_grids_svm = {'svc': {'object': svm.SVC(),
    'C': [10000, 100000, 1000000],
    'kernel': ['rbf', 'poly']
    }
}

mt_param_grids_list_svm = make_param_grids(mt_pipeline_steps_svm,
    ↪mt_all_param_grids_svm)
mt_pipe_svm = pipeline.make_pipeline(chosen_os, chosen_us, svm.SVC())
mt_grid_svm = model_selection.GridSearchCV(mt_pipe_svm, param_grid =
    ↪mt_param_grids_list_svm, scoring='accuracy', cv=5, verbose=15)
mt_grid_svm.fit(X_mt_numerical, y_mt)
```

The result is:

```
SVC(C=100000)
```

So now we have our 2 models and we can finally get to trying them on the test set!

```
[ ]: # Prepare estimators:
xgbc = XGBClassifier(alpha=0, colsample_bytree=0.8, eval_metric='mlogloss',
    ↪reg_lambda=0,
    max_depth=12, min_child_weight=3, n_estimators=250,
    objective='multi:softmax', random_state=1337, subsample=1.0,
    tree_method='gpu_hist', use_label_encoder=False)
rfc = ensemble.RandomForestClassifier(criterion='entropy', n_estimators=200,
    ↪random_state=1337)
svmc = svm.SVC(C=100000)

estimators_dict = {
    'xgbc': pipeline.make_pipeline(chosen_os, chosen_us, xgbc),
    'rfc': pipeline.make_pipeline(chosen_os, chosen_us, rfc),
    'svmc': pipeline.make_pipeline(chosen_os, chosen_us, svmc)
}

# Prepare test_data:
```

```

X_train_numerical = get_final_matrix(X_train, feature_strings_dict, brands_ohe,
    ↪corr_columns_to_drop, chosen_imputer)

# Fit
for name, estimator in estimators_dict.items():
    print("Fitting {}".format(name))
    estimator.fit(X_train_numerical, y_train)
    print("Finished fitting {}".format(name))

```

```

Fitting xgbc
Finished fitting xgbc
Fitting rfc
Finished fitting rfc
Fitting svmc
Finished fitting svmc

```

Let us compute the score on the test set:

```

[ ]: accs = {}
X_test_numerical = get_final_matrix(X_test, feature_strings_dict, brands_ohe,
    ↪corr_columns_to_drop, chosen_imputer)
for name, estimator in estimators_dict.items():
    print("Predicting {}".format(name))
    y_pred = estimator.predict(X_test_numerical)
    accs[name] = metrics.accuracy_score(y_test, y_pred)
    print("Finished predicting {}".format(name))
print(accs)

```

```

Predicting xgbc
Finished predicting xgbc
Predicting rfc
Finished predicting rfc
Predicting svmc
Finished predicting svmc
{'xgbc': 0.9310344827586207, 'rfc': 0.9266257282317745, 'svmc':
0.9044245000787278}

```

Not too bad. Now let us train our classifiers on all the data and make predictions for the test csv.

```

[ ]: X_numerical = get_final_matrix(X, feature_strings_dict, brands_ohe,
    ↪corr_columns_to_drop, chosen_imputer)
final_estimators_dict = {
    'xgbc': pipeline.make_pipeline(chosen_os, chosen_us, xgbc),
    'rfc': pipeline.make_pipeline(chosen_os, chosen_us, rfc),
    'svmc': pipeline.make_pipeline(chosen_os, chosen_us, svmc)
}
for name, estimator in final_estimators_dict.items():
    print("Fitting {}".format(name))

```

```
estimator.fit(X_numerical, y)
print("Finished fitting {}".format(name))
```

```
Fitting xgbc
Finished fitting xgbc
Fitting rfc
Finished fitting rfc
Fitting svmc
Finished fitting svmc
```

```
[ ]: test = pd.read_csv('drive/MyDrive/Final_Project-orav94/data/food_test.csv')
joined_test = pd.merge(test, nutrients, how='left', on='idx').
↳drop(['nutrient_id', 'unit_name'], axis=1)
pivoted_test = pd.pivot_table(
    joined_test,
    values='amount',
    index='idx',
    columns='name')

test_w_nutrients = pd.merge(test, pivoted_test, how='left', on='idx')
test_ids = test_w_nutrients[['idx']]
excl = ['alcohol_ethyl_g', 'energy_kj', 'folic_acid_ug', 'molybdenum_mo_ug',
↳'starch_g']
test_col_drop = [e for e in columns_to_drop[:len(columns_to_drop)-2] if e not
↳in excl]
test_final = clean_text_values(test_w_nutrients.drop(test_col_drop, axis=1)).
↳set_index('idx')
test_numerical = get_final_matrix(test_final, feature_strings_dict, brands_ohe,
↳corr_columns_to_drop, chosen_imputer)

preds = {}
for name, estimator in final_estimators_dict.items():
    print("Predicting {}".format(name))
    preds[name] = estimator.predict(test_numerical)
    print("Finished predicting {}".format(name))
```

```
Predicting xgbc
Finished predicting xgbc
Predicting rfc
Finished predicting rfc
Predicting svmc
Finished predicting svmc
```

```
[ ]: def decode(n):
    decode_dict = {
        0: 'cakes_cupcakes_snack_cakes', 1: 'candy', 2: 'chocolate', 3:
↳'cookies_biscuits',
```

```

        4: 'popcorn_peanuts_seeds_related_snacks', 5: 'chips_pretzels_snacks'
    }
    return decode_dict[n]

preds_dfs = {name: pd.concat([test_ids, pd.DataFrame(pred)], axis=1) for name, pred in preds.items()}
for name in preds_dfs.keys():
    preds_dfs[name] = preds_dfs[name].rename(columns={0: 'category_enc'})
    preds_dfs[name]['pred'] = preds_dfs[name]['category_enc'].apply(decode)
    preds_dfs[name] = preds_dfs[name].drop(['category_enc'], axis=1)
    preds_dfs[name].to_csv(path_or_buf='data/pred_{}.csv'.format(name),
        columns=['idx', 'pred'], index=False)

```